

```
In [2]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import MinMaxScaler
from sklearn.preprocessing import StandardScaler
from sklearn.preprocessing import LabelEncoder
from sklearn.metrics import precision_score, recall_score, accuracy_score, confusion_matrix
from sklearn.svm import SVC
from sklearn.linear_model import LogisticRegression
from sklearn.neighbors import KNeighborsClassifier
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import BaggingClassifier, AdaBoostClassifier, RandomForestClassifier, GradientBoostingClassifier
```

```
In [3]: df=pd.read_csv("C:/Users/hp/Downloads/loan_data.csv")
```

```
In [4]: df.head()
```

```
Out[4]:
```

	Loan_ID	Gender	Married	Dependents	Education	Self_Employed	ApplicantIncome	CoapplicantIncome	LoanAmount	Loan_Amount_Term	Credit_Hist
0	LP001003	Male	Yes	1	Graduate	No	4583	1508.0	128.0	360.0	
1	LP001005	Male	Yes	0	Graduate	Yes	3000	0.0	66.0	360.0	
2	LP001006	Male	Yes	0	Not Graduate	No	2583	2358.0	120.0	360.0	
3	LP001008	Male	No	0	Graduate	No	6000	0.0	141.0	360.0	
4	LP001013	Male	Yes	0	Not Graduate	No	2333	1516.0	95.0	360.0	

```
In [5]: df.columns
```

```
Out[5]: Index(['Loan_ID', 'Gender', 'Married', 'Dependents', 'Education',
'Self_Employed', 'ApplicantIncome', 'CoapplicantIncome', 'LoanAmount',
'Loan_Amount_Term', 'Credit_History', 'Property_Area', 'Loan_Status'],
dtype='object')
```

```
In [6]: df.info() # information of dataset
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 381 entries, 0 to 380
Data columns (total 13 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Loan_ID                381 non-null   object
1   Gender                 376 non-null   object
2   Married                381 non-null   object
3   Dependents             373 non-null   object
4   Education              381 non-null   object
5   Self_Employed          360 non-null   object
6   ApplicantIncome        381 non-null   int64
7   CoapplicantIncome      381 non-null   float64
8   LoanAmount             381 non-null   float64
9   Loan_Amount_Term       370 non-null   float64
10  Credit_History          351 non-null   float64
11  Property_Area           381 non-null   object
12  Loan_Status             381 non-null   object
dtypes: float64(4), int64(1), object(8)
memory usage: 38.8+ KB
```

```
In [7]: df.dtypes # Types of DataTypes present in DataSet
```

```
Out[7]: Loan_ID                object
Gender                 object
Married                object
Dependents             object
Education              object
Self_Employed          object
ApplicantIncome        int64
CoapplicantIncome      float64
LoanAmount             float64
Loan_Amount_Term       float64
Credit_History         float64
Property_Area           object
Loan_Status             object
dtype: object
```

```
In [8]: df.shape # Shape of Dataset i.e Number of Rows and Columns Present in Dataset
```

```
Out[8]: (381, 13)
```



McAfee WebAdvisor

Your download's being scanned.
We'll let you know if there's an issue.

```
In [9]: df.size # Size of DataSet i.e Total number of Elements Present in Dataset
```

```
Out[9]: 4953
```

```
In [10]: df.isnull().sum() # Checking Null Values present in dataset
```

```
Out[10]: Loan_ID      0
Gender      5
Married     0
Dependents  8
Education   0
Self_Employed  21
ApplicantIncome  0
CoapplicantIncome  0
LoanAmount   0
Loan_Amount_Term  11
Credit_History  30
Property_Area  0
Loan_Status  0
dtype: int64
```

```
In [11]: df1=df.dropna() # Dropping Null Values present in Dataset
```

```
In [12]: df1.isnull().sum()
```

```
Out[12]: Loan_ID      0
Gender      0
Married     0
Dependents  0
Education   0
Self_Employed  0
ApplicantIncome  0
CoapplicantIncome  0
LoanAmount   0
Loan_Amount_Term  0
Credit_History  0
Property_Area  0
Loan_Status  0
dtype: int64
```

```
In [13]: df1.dtypes
```

```
Out[13]: Loan_ID      object
Gender      object
Married     object
Dependents  object
Education   object
Self_Employed  object
ApplicantIncome  int64
CoapplicantIncome  float64
LoanAmount      float64
Loan_Amount_Term  float64
Credit_History  float64
Property_Area    object
Loan_Status      object
dtype: object
```

Performing LabelEncoding On DataSet

```
In [14]: le=LabelEncoder() # To convert categorical data into numerical data for model building
```



McAfee WebAdvisor



Your download's being scanned.
We'll let you know if there's an issue.

```
In [15]: df1["Loan_ID"]=le.fit_transform(df1["Loan_ID"])
df1["Gender"]=le.fit_transform(df1["Gender"])
df1["Married"]=le.fit_transform(df1["Married"])
df1["Dependents"]=le.fit_transform(df1["Dependents"])
df1["Education"]=le.fit_transform(df1["Education"])
df1["Self_Employed"]=le.fit_transform(df1["Self_Employed"])
df1["Property_Area"]=le.fit_transform(df1["Property_Area"])
df1["Loan_Status"]=le.fit_transform(df1["Loan_Status"])

C:\Users\hp\AppData\Local\Temp\ipykernel_12844\4008765204.py:1: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user\_guide/indexing.html#returning-a-view-versus-a-copy
df1["Loan_ID"]=le.fit_transform(df1["Loan_ID"])
C:\Users\hp\AppData\Local\Temp\ipykernel_12844\4008765204.py:2: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user\_guide/indexing.html#returning-a-view-versus-a-copy
df1["Gender"]=le.fit_transform(df1["Gender"])
C:\Users\hp\AppData\Local\Temp\ipykernel_12844\4008765204.py:3: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user\_guide/indexing.html#returning-a-view-versus-a-copy
df1["Married"]=le.fit_transform(df1["Married"])
C:\Users\hp\AppData\Local\Temp\ipykernel_12844\4008765204.py:4: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user\_guide/indexing.html#returning-a-view-versus-a-copy
df1["Dependents"]=le.fit_transform(df1["Dependents"])
C:\Users\hp\AppData\Local\Temp\ipykernel_12844\4008765204.py:5: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user\_guide/indexing.html#returning-a-view-versus-a-copy
df1["Education"]=le.fit_transform(df1["Education"])
C:\Users\hp\AppData\Local\Temp\ipykernel_12844\4008765204.py:6: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user\_guide/indexing.html#returning-a-view-versus-a-copy
df1["Self_Employed"]=le.fit_transform(df1["Self_Employed"])
C:\Users\hp\AppData\Local\Temp\ipykernel_12844\4008765204.py:7: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user\_guide/indexing.html#returning-a-view-versus-a-copy
df1["Property_Area"]=le.fit_transform(df1["Property_Area"])
C:\Users\hp\AppData\Local\Temp\ipykernel_12844\4008765204.py:8: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user\_guide/indexing.html#returning-a-view-versus-a-copy
df1["Loan_Status"]=le.fit_transform(df1["Loan_Status"])
```

```
In [16]: df1.head()
```

```
Out[16]:
```

	Loan_ID	Gender	Married	Dependents	Education	Self_Employed	ApplicantIncome	CoapplicantIncome	LoanAmount	Loan_Amount_Term	Credit_Histo
0	0	1	1	1	0	0	4583	1508.0	128.0	360.0	1
1	1	1	1	0	0	1	3000	0.0	66.0	360.0	1
2	2	1	1	0	1	0	2583	2358.0	120.0	360.0	1
3	3	1	0	0	0	0	6000	0.0	141.0	360.0	1
4	4	1	1	0	1	0	2333	1516.0	95.0	360.0	1



McAfee WebAdvisor



Your download's being scanned.
We'll let you know if there's an issue.

```
In [17]: df1.describe().T      # Statistical Summary Of Dataset
```

```
Out[17]:
```

	count	mean	std	min	25%	50%	75%	max
Loan_ID	308.0	153.500000	89.056162	0.0	76.75	153.5	230.25	307.0
Gender	308.0	0.795455	0.404025	0.0	1.00	1.0	1.00	1.0
Married	308.0	0.600649	0.490562	0.0	0.00	1.0	1.00	1.0
Dependents	308.0	0.678571	0.997029	0.0	0.00	0.0	1.00	3.0
Education	308.0	0.256494	0.437408	0.0	0.00	0.0	1.00	1.0
Self_Employed	308.0	0.090909	0.287948	0.0	0.00	0.0	0.00	1.0
ApplicantIncome	308.0	3599.126623	1462.359612	150.0	2568.75	3329.5	4291.00	9703.0
CoapplicantIncome	308.0	1278.434805	2520.961308	0.0	0.00	871.5	1953.50	33837.0
LoanAmount	308.0	104.623377	29.382256	9.0	89.75	110.0	128.00	150.0
Loan_Amount_Term	308.0	341.181818	68.246006	36.0	360.00	360.0	360.00	480.0
Credit_History	308.0	0.853896	0.353785	0.0	1.00	1.0	1.00	1.0
Property_Area	308.0	1.042208	0.775125	0.0	0.00	1.0	2.00	2.0
Loan_Status	308.0	0.711039	0.454017	0.0	0.00	1.0	1.00	1.0

Changing DataTypes of Data in Dataset

```
In [18]: df1["CoapplicantIncome"]=df1["CoapplicantIncome"].astype("int")
df1["LoanAmount"]=df1["LoanAmount"].astype("int")
df1["Loan_Amount_Term"]=df1["Loan_Amount_Term"].astype("int")
df1["Credit_History"]=df1["Credit_History"].astype("int")
```

C:\Users\hp\AppData\Local\Temp\ipykernel_12844\1689938164.py:1: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy (https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)
df1["CoapplicantIncome"]=df1["CoapplicantIncome"].astype("int")

C:\Users\hp\AppData\Local\Temp\ipykernel_12844\1689938164.py:2: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy (https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)
df1["LoanAmount"]=df1["LoanAmount"].astype("int")

C:\Users\hp\AppData\Local\Temp\ipykernel_12844\1689938164.py:3: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy (https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)
df1["Loan_Amount_Term"]=df1["Loan_Amount_Term"].astype("int")

C:\Users\hp\AppData\Local\Temp\ipykernel_12844\1689938164.py:4: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy (https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)
df1["Credit_History"]=df1["Credit_History"].astype("int")



McAfee WebAdvisor

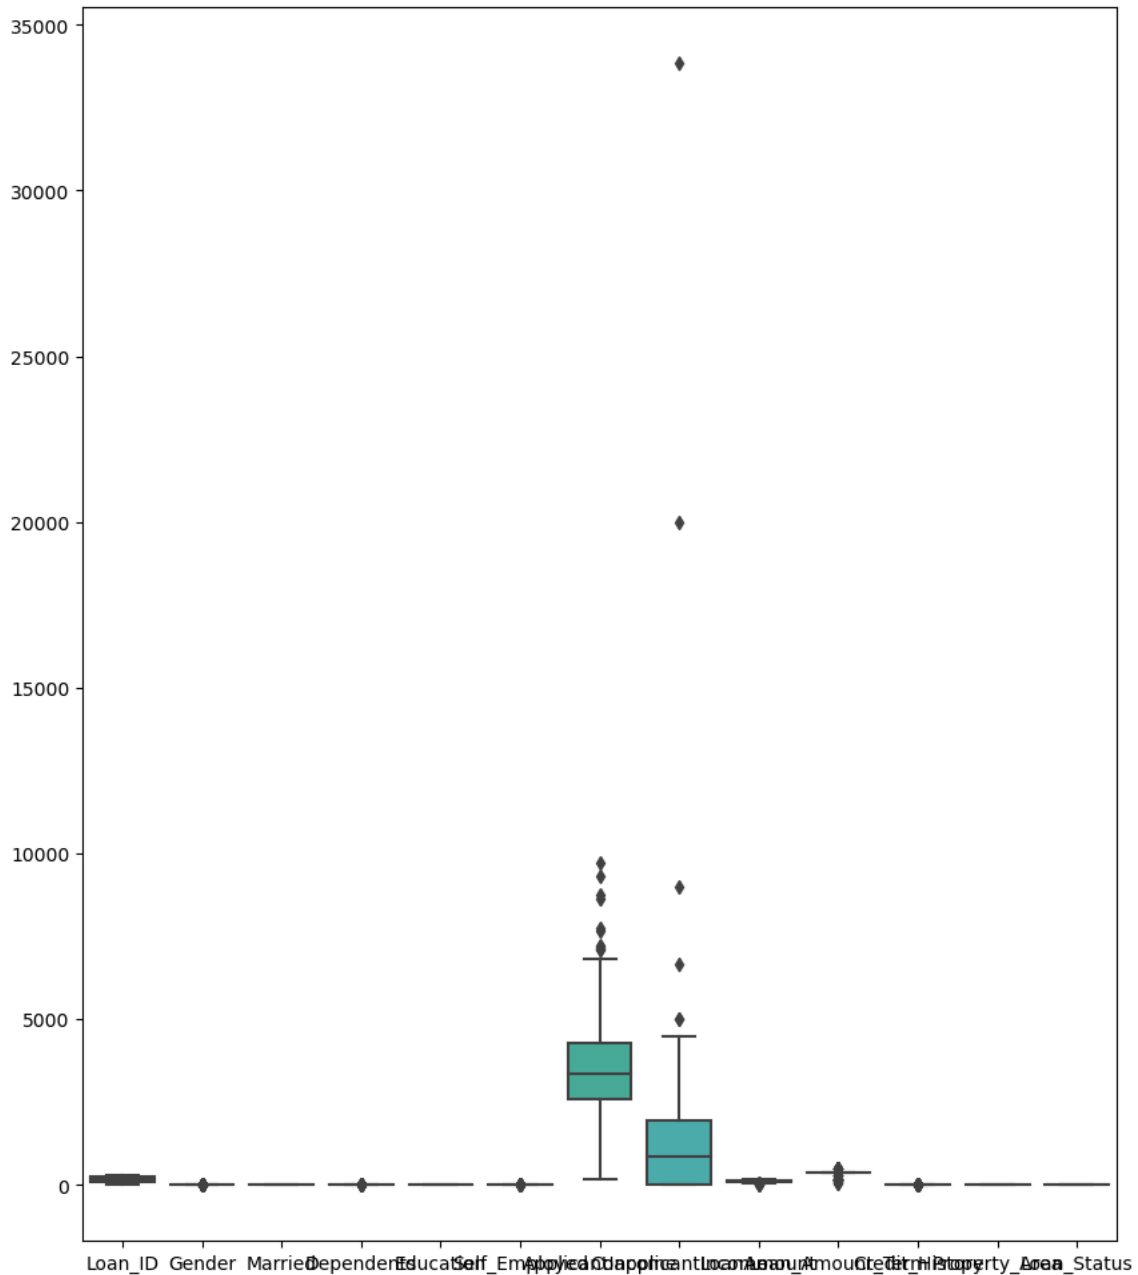


Your download's being scanned.
We'll let you know if there's an issue.

Removing Outliers present in Dataset

```
In [19]: plt.figure(figsize=(10,12))          # Ploting Boxplot to detect outliers
sns.boxplot(data=df1)
```

Out[19]: <Axes: >



Steps to Remove Outliers

```
In [20]: Q1=df1.quantile(q=0.25)          # finding Q1 value
Q3=df1.quantile(q=0.75)                  # Finding Q3 value
IQR=Q3-Q1                                # Finding IQR Value i.e.(InterQuantileRange)
upper=Q3+(1.5*IQR)                       # to detect upper outliers
lower=Q1-(1.5*IQR)                       # to detect Lower outliers
```

```
In [21]: df2=df1[~((df1>upper)|(df1<lower))]
```




McAfee WebAdvisor



Your download's being scanned.
We'll let you know if there's an issue.

```
In [22]: plt.figure(figsize=(10,12)) # Boxplot after removing Outliers
sns.boxplot(df2)
```

Out[22]: <Axes: >



```
In [23]: df2.isnull().sum() # After removing outlier we get some nan values
```

```
Out[23]: Loan_ID      0
Gender      63
Married     0
Dependents  24
Education   0
Self_Employed  28
ApplicantIncome  11
CoapplicantIncome  6
LoanAmount   7
Loan_Amount_Term  49
Credit_History  45
Property_Area  0
Loan_Status  0
dtype: int64
```

```
In [24]: df3=df2.dropna() # removing that nan values
```




McAfee WebAdvisor



Your download's being scanned.
We'll let you know if there's an issue.

HeatMap To Show Correlation between Data

```
In [25]: plt.figure(figsize=(12,10))
sns.heatmap(df2.corr(),annot=True)
plt.show()
```



```
In [26]: print(df1["Loan_Status"].unique())
print(df1["Loan_Status"].value_counts(normalize=True)*100)
```

```
[0 1]
Loan_Status
1    71.103896
0    28.896104
Name: proportion, dtype: float64
```




McAfee WebAdvisor



Your download's being scanned.
We'll let you know if there's an issue.

```
In [27]: plt.pie(df2["Loan_Status"].value_counts(normalize=True)*100,labels=["YES","NO"])
```

```
Out[27]: ([<matplotlib.patches.Wedge at 0x294fad145d0>,
<matplotlib.patches.Wedge at 0x294fad2e7d0>],
[Text(-0.6770310906241628, 0.866965340903693, 'YES'),
Text(0.6770310906241627, -0.8669653409036931, 'NO')])
```



Model Building for DataSet

```
In [28]: x=df3.drop(["Loan_Status"],axis=1)
y=df3["Loan_Status"]
```

```
In [29]: x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3,random_state=1) # splitting data into x and y
```

```
In [30]: print(x_train.shape)
print(x_test.shape)
print(y_train.shape)
print(y_test.shape)
```

```
(99, 12)
(43, 12)
(99,)
(43,)
```

LOGISTIC REGRESSION ALGORITHM

```
In [31]: le=LogisticRegression()
le.fit(x_train,y_train)
```

```
Out[31]: LogisticRegression()
```

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.
On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

```
In [32]: y_true,y_pred=y_test,le.predict(x_test)
print(le.score(x_train,y_train)*100)
print(le.score(x_test,y_test)*100)
```

```
86.86868686868688
88.37209302325581
```

```
In [33]: print(precision_score(y_true,y_pred)*100)
print(recall_score(y_true,y_pred)*100)
print(accuracy_score(y_true,y_pred)*100)
```

```
88.37209302325581
100.0
88.37209302325581
```

RANDOM FOREST CLASSIFIER ALGORITHM

```
In [34]: rf=RandomForestClassifier(n_estimators=6,random_state=1)
rf.fit(x_train,y_train)
```

```
Out[34]: RandomForestClassifier(n_estimators=6, random_state=1)
```

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.
On GitHub, the HTML representation is unable to render, please try loading this page



McAfee WebAdvisor



Your download's being scanned.
We'll let you know if there's an issue.


```
In [35]: y_true,y_pred=y_test,rf.predict(x_test)
print(rf.score(x_train,y_train)*100)
print(rf.score(x_test,y_test)*100)
```

```
95.95959595959596
83.72093023255815
```

```
In [36]: print(precision_score(y_true,y_pred)*100)
print(recall_score(y_true,y_pred)*100)
print(accuracy_score(y_true,y_pred)*100)
```

```
87.8048780487805
94.73684210526315
83.72093023255815
```

DECISION TREE CLASSIFIER ALGORITHM

```
In [37]: dt=DecisionTreeClassifier(criterion="gini",max_depth=4,random_state=1)
dt.fit(x_train,y_train)
```

Out[37]: DecisionTreeClassifier(max_depth=4, random_state=1)

**In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.
On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.**

```
In [38]: y_true,y_pred=y_test,dt.predict(x_test)
print(dt.score(x_train,y_train)*100)
print(dt.score(x_test,y_test)*100)
```

```
94.94949494949495
88.37209302325581
```

```
In [39]: print(precision_score(y_true,y_pred)*100)
print(recall_score(y_true,y_pred)*100)
print(accuracy_score(y_true,y_pred)*100)
```

```
90.2439024390244
97.36842105263158
88.37209302325581
```

GRADIENT BOOSTING CLASSIFIER ALGORITHM

```
In [40]: gb=GradientBoostingClassifier(n_estimators=20)
gb.fit(x_train,y_train)
```

Out[40]: GradientBoostingClassifier(n_estimators=20)

**In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.
On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.**

```
In [41]: y_true,y_pred=y_test,gb.predict(x_test)
print(dt.score(x_train,y_train)*100)
print(dt.score(x_test,y_test)*100)
```

```
94.94949494949495
88.37209302325581
```

```
In [42]: print(precision_score(y_true,y_pred)*100)
print(recall_score(y_true,y_pred)*100)
print(accuracy_score(y_true,y_pred).*100)
```

```
90.47619047619048
100.0
90.69767441860465
```

BAGGING CLASSIFIER ALGORITHM

```
In [43]: bg=BaggingClassifier(n_estimators=20)
bg.fit(x_train,y_train)
```

Out[43]: BaggingClassifier(n_estimators=20)

**In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.
On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.**

```
In [44]: y_true,y_pred=y_test,bg.predict(x_test)
print(dt.score(x_train,y_train)*100)
print(dt.score(x_test,y_test)*100)
```

```
94.94949494949495
88.37209302325581
```



McAfee WebAdvisor



Your download's being scanned.
We'll let you know if there's an issue.

```
In [45]: print(precision_score(y_true,y_pred)*100)
print(recall_score(y_true,y_pred)*100)
print(accuracy_score(y_true,y_pred)*100)
```

```
90.47619047619048
100.0
90.69767441860465
```

ADABOOST CLASSIFIER ALGORITHM

```
In [46]: ad=AdaBoostClassifier(n_estimators=20,estimator=dt,random_state=1)
ad.fit(x_train,y_train)
```

```
Out[46]: AdaBoostClassifier(estimator=DecisionTreeClassifier(max_depth=4,
random_state=1),
n_estimators=20, random_state=1)
```

**In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.
On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.**

```
In [47]: y_true,y_pred=y_test,ad.predict(x_test)
print(dt.score(x_train,y_train)*100)
print(dt.score(x_test,y_test)*100)
```

```
94.94949494949495
88.37209302325581
```

```
In [48]: print(precision_score(y_true,y_pred)*100)
print(recall_score(y_true,y_pred)*100)
print(accuracy_score(y_true,y_pred)*100)
```

```
90.47619047619048
100.0
90.69767441860465
```

KNeighbors CLASSIFIER ALGORITHM

```
In [49]: kn=KNeighborsClassifier(weights="distance")
kn.fit(x_train,y_train)
```

```
Out[49]: KNeighborsClassifier(weights='distance')
```

**In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.
On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.**

```
In [50]: y_true,y_pred=y_test,kn.predict(x_test)
print(dt.score(x_train,y_train)*100)
print(dt.score(x_test,y_test)*100)
```

```
94.94949494949495
88.37209302325581
```

```
In [51]: print(precision_score(y_true,y_pred)*100)
print(recall_score(y_true,y_pred)*100)
print(accuracy_score(y_true,y_pred)*100)
```

```
87.5
92.10526315789474
81.3953488372093
```

SVC (SUPPORT VECTOR CLASSIFIER) ALGORITHM

```
In [52]: svc=SVC(C=1.0,kernel="linear")
svc.fit(x_train,y_train)
```

```
Out[52]: SVC(kernel='linear')
```

**In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.
On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.**

```
In [53]: y_true,y_pred=y_test,svc.predict(x_test)
print(dt.score(x_train,y_train)*100)
print(dt.score(x_test,y_test)*100)
```

```
94.94949494949495
88.37209302325581
```

```
In [54]: print(precision_score(y_true,y_pred)*100)
print(recall_score(y_true,y_pred)*100)
print(accuracy_score(y_true,y_pred)*100)
```

```
88.37209302325581
100.0
88.37209302325581
```



McAfee WebAdvisor



Your download's being scanned.
We'll let you know if there's an issue.

In []: *#The best fit model for this bank loan data set are Gradient boosting algorithm, bagging classifier algorithm, adaboost clas*

In []:



McAfee WebAdvisor



Your download's being scanned.
We'll let you know if there's an issue.